

# Some Deep Learning Cooking Recipes in TensorFlow

Thursday 7<sup>th</sup> March, 2019



**Matthia Sabatelli<sup>1</sup>**

<sup>1</sup>Montefiore Institute, Department of Electrical Engineering and Computer Science, Université de Liège, Belgium

1. Back in the days...
2. Tensorflow
3. Keras

Back in the days...

## Back in the days...

```
1 def calculate_loss(model):
2     num_examples = len(X)
3
4     W1, b1, W2, b2 = model['W1'], model['b1'],
5         model['W2'], model['b2']
6
7     z1 = X.dot(W1) + b1
8     a1 = np.tanh(z1)
9     z2 = a1.dot(W2) + b2
10
11     exp_scores = np.exp(z2)
12     probs = exp_scores / np.sum(exp_scores, axis=1,
13         keepdims=True)
14     corect_logprobs =
15         -np.log(probs[range(num_examples), y])
16
17     data_loss = np.sum(corect_logprobs)
18     data_loss += reg_lambda / 2 *
19         (np.sum(np.square(W1)) +
20          np.sum(np.square(W2)))
```

## Back in the days...

```
1 public class MLP implements Cloneable
2 {
3     protected double fLearningRate = 0.001;
4     protected Layer[] fLayers;
5     double dVal=0;
6     public ArrayList<Integer> checkArray = new
        ArrayList<Integer>();
7
8     public MLP(int[] layers, double learningRate){
9         fLearningRate = learningRate;
10        fLayers = new Layer[layers.length];
11        for(int i = 0; i < layers.length; i++){
12            if(i != 0){
13                fLayers[i] = new Layer(layers[i],
                    layers[i - 1]);
14            }
15            else{
16                ...
```

## Back in the days...

```
1 import numpy
2 import theano as T
3
4 w_1 = T.shared(rng.randn(784, 300), name='w1')
5 b_1 = T.shared(numpy.zeros((300,)), name='b1')
6
7 from theano.tensor.nnet import sigmoid
8
9 p_1 = sigmoid(-T.dot(sigmoid(-T.dot(x, w_1)-b_1),
10                        w_2)-b_2)
11 xent = -y*T.log(p_1) - (1-y)*T.log(1-p_1)
12 cost = xent.mean() + 0.01 * (w_2** 2).sum()
13
14 gw_1, gb_1, gw_2, gb_2 = T.grad(cost, [w_1, b_1, w_2,
15                                         b_2])
16
17 train = T.function(inputs = [x, y], outputs =
18                     [prediction, xent], updates = {w_1 :
19                                                     w_1-0.1*gw_1, b_1 : b_1-0.1*gb_1, ... })
20
21 predict = T.function(inputs=[x], outputs=prediction)
```

Tensorflow

## Some TF facts

- Developed by Google Brain and open-sourced in November 2015
- Not only limited to Deep Learning
- Supports many Python APIs (Keras, TF-slim, TF-Learn)
- Runs efficient C++ code in the backend
- Provides **Automatic Differentiation!**



## Install Anaconda and create a Virtual-Env

```
1 curl -O https://repo.continuum.io/archive/  
2     Anaconda3-5.0.1-Linux-x86_64.sh  
3  
4 bash Anaconda3-5.0.1-Linux-x86_64.sh  
5  
6 conda create --name my_tf_env python=3  
7  
8 source activate my_tf_env
```

## Install appropriate Tf version

```
1 conda install -c anaconda tensorflow  
2  
3 conda install -c anaconda tensorflow-gpu
```

## Check if GPUs are properly seen

```
1 import tensorflow as tf  
2  
3 sess = tf.Session(config=tf.ConfigProto(  
4     log_device_placement=True)
```

## If everything goes well...

```
1 name: GeForce GTX 1080 Ti
2 major: 6 minor: 1 memoryClockRate (GHz) 1.582
3 pciBusID 0000:82:00.0
4 Total memory: 10.92GiB
5 Free memory: 10.00GiB
6 2019-03-07 09:12:30.130236: I
   tensorflow/core/common_runtime/gpu/gpu_device.cc:976]
   DMA: 0
7 2019-03-07 09:12:30.130254: I
   tensorflow/core/common_runtime/gpu/gpu_device.cc:986]
   0: Y
8 2019-03-07 09:12:30.130266: I
   tensorflow/core/common_runtime/gpu/gpu_device.cc:1045]
   Creating TensorFlow device (/gpu:0) -> (device:
   0, name: GeForce GTX 1080 Ti, pci bus id:
   0000:82:00.0)
9 Device mapping:
10 /job:localhost/replica:0/task:0/gpu:0 -> device: 0,
   name: GeForce GTX 1080 Ti, pci bus id:
   0000:82:00.0
```

# Our First (Basic) Computation Graph

```
1 import tensorflow as tf
2
3 x = tf.Variable(3, name="x")
4 y = tf.Variable(4, name="y")
5 f = x*x*y+y+2
6
7 sess = tf.Session()
8 sess.run(x.initializer)
9 sess.run(y.initializer)
10
11 result = sess.run(f)
12
13 sess.close()
```

```
1 import tensorflow as tf
2
3 init = tf.global_variables_initializer()
4
5 with tf.Session() as sess:
6     init.run()
7     result = f.eval()
```

# Our First (Ugly) Neural Network

```
1 import tensorflow as tf
2
3 n_inputs = 28*28
4 n_hidden_1 = 300
5 n_hidden_2 = 100
6 n_outputs = 10
7
8 X = tf.placeholder(tf.float32, shape(None, n_inputs),
9                    name="X") # input layer
10 y = tf.placeholder(tf.int64, shape(None),
11                   name="y") # output layer
```

- We want 2 hidden layers
- Introduce non-linearity

# Our First (Ugly) Neural Network

```
1 def neuron_layer(X, n_neurons, name):
2     with tf.name_scope(name):
3         n_inputs = int(X.get_shape()[1])
4         stddev = 2 / np.sqrt(n_inputs + n_neurons)
5         init = tf.truncated_normal((n_inputs,
6                                     n_neurons), stddev=stddev)
7
8         W = tf.Variable(init, name="weight_matrix")
9         b = tf.Variable(tf.ones([n_neurons]),
10                          name="bias")
11         Z = tf.matmul(X,W)+b
12
13         return tf.nn.relu(Z)
```

```
1 with tf.name_scope("my_net"):
2     hidden_1 = neuron_layer(X, n_hidden_1, name =
3                             "hidden_1")
4     hidden_2 = neuron_layer(hidden_1, n_hidden_2,
5                             name = "hidden_2")
6     logits = neuron_layer(hidden_2, n_outputs, name
7                             = "outputs")
```

# Who wants to program a layer???!?

Tf-Layers to the rescue!

```
1 with tf.name_scope("my_net"):  
2     hidden_1 = tf.layers.dense(X, n_hidden_1, name =  
3         "hidden_1", activation = tf.nn.relu)  
4     hidden_2 = tf.layers.dense(hidden_1, n_hidden_2,  
5         name = "hidden_2", activation = tf.nn.relu)  
6     logits = tf.layers.dense(hidden_2, n_outputs,  
7         name = "outputs") # output before final  
8                             activation function
```

However we still need to train our graph somehow:

- Cost function
- Optimizer

## Our Loss

```
1 with tf.name_scope("loss"):
2     entropy =
        tf.nn.sparse_softmax_cross_entropy_with_logits(labels=
        logits=logits)
```

## Our Optimizer

```
1 with tf.name_scope("train"):
2     optimizer =
        tf.train.GradientDescentOptimizer(0.0001)
3     training_op = optimizer.minimize(loss)
```

## Our Evaluation Metric

```
1 with tf.name_scope("eval"):
2     correct = tf.nn.in_top_k(logits, y, 1) # Does
        the highest logit match with the class?
3     accuracy = tf.reduce_mean(tf.cast(correct,
        tf.float32))
```

What we have done so far

- We created the placeholders for the inputs and targets
- Defined the layers of the network
- Defined an objective function to minimize
- Initialized the optimizer
- Defined a performance measure

⇒ End of the **Construction-Phase** of the graph

⇒ Let's move to the **Execution-Phase!**




## Training Time

```
1 with tf.Session as sess:
2     init.run()
3     for epoch in range(n_epochs):
4         for iteration in
5             range(mnist.train.num_examples //
6                 batch_size):
7                 X_batch, y_batch =
8                     mnist.train.next_batch(batch_size)
9                     sess.run(training_op, feed_dict={X:
10                        X_batch, y: y_batch})
11
12     training_acc = accuracy.eval(feed_dict={X:
13        X_batch, y: y_batch})
14     val_acc = accuracy.eval(feed_dict={X:
15        mnist.validation_images, y:
16        mnist.validation_labels})
```

## Testing Time

```
1 with tf.Session as sess:  
2     saver.restore(sess, "./my_path/model.ckpt")  
3     Z =  
4         logits.eval(feed_dict={X:mnist.testing_images})  
5     y_pred = np.argmax(Z, axis = 1)
```

# Activation Functions



- `tf.nn.relu`
- `tf.nn.elu`
- `tf.nn.sigmoid`
- `tf.nn.tanh`
- `tf.nn.linear`

Figure: Graphical representation of some potential activation functions

```
1 def leaky_relu(z, name=None):  
2     return tf.maximum(0.01 * z, z, name=name)  
3  
4 hidden_layer = tf.layers.dense(X, n_hidden1,  
    activation=leaky_relu, name="hidden1")
```

- `tf.nn.mean_squared_error()`
- `tf.nn.sigmoid_cross_entropy()`
- `tf.nn.softmax_cross_entropy()`
- `tf.nn.cosine_distance()`
- `tf.nn.hinge_loss()`

## Custom Loss Function

```
1 self.target_q_t = tf.placeholder('float32', [None],  
    name='target_q_t')  
2 self.action = tf.placeholder('int64', [None],  
    name='action')  
3  
4 action_one_hot = tf.one_hot(self.action,  
    self.env.action_size, 1.0, 0.0,  
    name='action_one_hot')  
5  
6 q_acted = tf.reduce_sum(self.q * action_one_hot,  
    reduction_indices=1, name='q_acted')  
7  
8 self.delta = self.target_q_t - q_acted  
9  
10 self.loss = tf.reduce_mean(clipped_error(self.delta),  
    name='loss')
```

- `tf.train.GradientDescent()`
- `tf.train.AdadeltaOptimizer()`
- `tf.train.AdamOptimizer()`
- `tf.train.RMSPropOptimizer()`
- `tf.train.MomentumOptimizer()`

⇒ It is easy to change the default parameters and to minimize a custom objective function!

```
1 optim = tf.train.RMSPropOptimizer(  
2     self.learning_rate_op, momentum=0.95,  
    epsilon=0.01).minimize(self.loss)
```

# Avoiding Overfitting through Regularization

```
1 my_regularized_layer = partial(tf.layers.dense,  
    activation=tf.nn.relu,  
    kernel_regularizer=tf.contrib.layers.l2_regularizer(0.0005))  
2  
3 with tf.name_scope("regularized_dnn"):  
4     hidden_1 = my_regularized_layer(X, 100,  
        name="reg_hidden1")  
5     hidden_2 = my_regularized_layer(hidden_1, 100,  
        name="reg_hidden2")  
6     logits = my_regularized_layer(hidden_2,  
        n_outputs, activation=None,  
        name="output_layer")
```

**Remember to add the regularized losses to your overall loss**

```
1 reg_losses =  
    tf.get_collection(tf.GraphKeys.REGULARIZATION_LOSSES)  
2 loss = tf.add_n([base_loss] + reg_losses, name  
    ="regularized_loss")
```

## Our own regularizer

Max-Norm Regularization:  $\sum_k \sum_i \sum_j (w_{i,j}^k)^2 \leq r$

```
1 def max_norm_regularizer(r, axes=1, name="max_norm",  
    collection="max_norm"):  
2  
3     def max_norm(w):  
4         clipped = tf.clip_by_norm(w, clip_norm = r,  
            axes = axes)  
5         clip_weights = tf.assign(w, clipped, name =  
            name)  
6         tf.add_to_collection(collection,  
            clip_weights)  
7  
8         return None      # no regularization loss is  
            required this time  
9  
10    return max_norm  
11  
12 hidden = tf.layers.dense(X, 100, activation =  
    tf.nn.relu, kernel_regularizer =  
    max_norm_regularizer, name = "hidden1")
```



# Our own regularizer

```
1 with tf.name_scope("my_reg_nn"):
2     hidden = tf.layers.dense(X, 100, activation =
        tf.nn.relu, kernel_regularizer =
        max_norm_regularizer, name = "hidden1")
```

# Dropout

```
1 training = tf.placeholder_with_default(True,  
    shape=(), "training")  
2  
3 dropout_rate = 0.2 # Reminder == 1 - keep_prob  
4 X_drop = tf.layers.dropout(X, dropout_rate,  
    training=training)  
5  
6 with tf.name_scope("dropout_nn"):  
7     hidden_1 = tf.layers.dense(X, 100,  
        activation=tf.nn.relu, name="hidden_1")  
8     hidden_1_dropout = tf.layers.dense(X_drop,  
        dropout_rate, training=training)
```

**Remember to set training to False at test time!**

Keras

## Sweet construction phase

- Super-simple
- Works only for single-input/output
- No more placeholders/scopes and sessions to initialize

```
1 import keras
2 from keras import layers
3
4 model = keras.Sequential()
5
6 model.add(layers.Dense(20, activation="relu",
7                        inpute_shape=(784,)))
7 model.add(layers.Dense(20, activation="relu"))
8 model.add(layers.Dense(10, activation="softmax"))
```

## Sweet evaluation phase

- No more annoying loops
- No more `init.run()` nor `sess.run()`

```
1 model.compile(loss='categorical_crossentropy',  
    optimizer='sgd', metrics=['accuracy'])  
2 model.fit(X, y, epochs = 10, batch_size = 32)  
3  
4 model.evaluate(X_test, y_test)
```

# The Functional API

- Similar to LEGO blocks
- Suitable for more complex architectures
- Allows multiple inputs/outputs

```
1 import keras
2 from keras import layers
3
4 input_frames = keras.Input(shape=(84, 84, 4))
5
6 x = layers.Conv2D(32, (8, 8),
7                   activation='relu')(input_frames)
8 x = layers.Conv2D(64, (3, 3), activation='relu')(x)
9 x = layers.Conv2D(32, (3, 3), activation='relu')(x)
10
11 flattened_representation = layers.Flatten()(x)
12 shared_representation = layers.Dense(512,
13                                       activation='relu')(flattened_representation)
```

# The Functional API

```
1 q_outputs = layers.Dense(3, activation =  
    'linear')(shared_representation)  
2 v_output = layers.Dense(1, activation =  
    'relu')(shared_representation)  
3  
4 model = keras.Model(inputs, outputs = [q_outputs,  
    v_output])
```

# Data Generators

```
1 import numpy as np
2
3 from keras import layers
4
5 X = np.load('my_huuuuuge_training_set.npy')
6 y = np.load('my_huuuuuuuuuuuge_labels.npy')
7
8 model = Sequential()
9 [...] # My fancy Neural Net
10 model.compile()
11
12 model.fit(X, y, epochs = 500, batch_size = 128)
```



## 1) We first define our Generator

```
1 from keras.preprocessing.image import  
    ImageDataGenerator  
2  
3 data_generator = ImageDataGenerator(  
4     rescale=1./255,  
5     horizontal_flip = True,  
6     vertical_flip = False,  
7     height_shift_range = 0.15,  
8     width_shift_range = 0.15,  
9     rotation_range = 5,  
10    shear_range = 0.01,  
11    fill_mode = 'nearest',  
12    zoom_range=0.25)
```

## 2) We yield appropriate batches of data

```
1 data_gen =  
    data_generator.flow_from_directory(IMAGES_PATH,  
2                                     target_size = (224,224),  
3                                     class_mode = 'categorical',  
4                                     batch_size = 32,  
5                                     shuffle=True,  
6                                     seed=7)  
7  
8 while True:  
9     data = data_gen.next()  
10    yield [data[0], data[1]]
```

- The `data_generator.*` method depends on your dataset structure!

## 3) We appropriately fit our model

```
1 step_size_train = data_gen.n // batch_size
2
3 model.fit_generator(generator = data_gen,
4                     steps_per_epoch = step_size_train,
5                     callbacks=[csv_logger, tbCallback]
6                     epochs=10)
```

- Similarly for the evaluation phase there is a `model.evaluate_generator` method

## Using ImageNet as a powerful Source-Domain

```
1 from keras.applications import ResNet50
2
3 base_model =
4     ResNet50(weights='imagenet', include_top=False)
5
6 x = base_model.output
7 x = GlobalAveragePooling2D()(x)
8 x = Dense(1024, activation='relu')(x)
9
10 preds = Dense(100, activation='softmax')(x)
11
12 model = Model(inputs = X, outputs = predictions)
```

- Which layers should use ImageNet's information?
- We can access this information as we would be dealing with lists

```
1 for layer in model.layers:  
2     layer.trainable=False  
3  
4 for layer in model.layers[20:]:  
5     layer.trainable=True
```

- One can explore which layers carry which features
- **Drastically reduce the amount of trainable parameters**

# Callbacks

```
1 model.fit_generator(generator=training_generator
2                     ,...,
3                     callbacks=[csv_logger, tbCallback]
4                     )
```

## Monitoring your results

```
1 csv_logger = CSVLogger(results_path + '/log_' +
    dataset_name + '.csv', append=True,
    separator='\t')
```

## Regularizing training

```
1 earlyStoppingCallback =
    keras.callbacks.EarlyStopping(monitor='val_loss',
    patience=5, verbose=0, mode='auto')
```

## Tensorboard is da bomb!

```
1 tbCallback =
    keras.callbacks.TensorBoard(log_dir=tensorboard_path,
    histogram_freq=0, write_graph=True,
    write_images=True)
```

## The CSVLogger Callback

- We can easily keep track of the training progress
- All results are neatly saved for us

1	epoch	tr_acc_1	tr_loss_1	val_loss_1	...
2	0	0.854	3.965	6.726	0.654
3	1	0.8729	3.215	4.792	0.772
4	2	0.8601	2.624	4.078	0.797

⇒ At the end of training we can easily process and plot our results!

- Géron, Aurélien. *Hands-on machine learning with Scikit-Learn and TensorFlow: concepts, tools, and techniques to build intelligent systems.* " O'Reilly Media, Inc.", 2017.
- Francois, Chollet. "Deep learning with Python." (2017).
- <https://www.tensorflow.org/tutorials>
- <https://github.com/fchollet/keras-resources>



Now it is up to you! Have fun :)